

Text Summarization using Deep Learning Model

Developed by

- [Islam Mohamed \(https://github.com/islammohamedd1\)](https://github.com/islammohamedd1)
- [Mostafa Yasser \(https://github.com/mostafayasser\)](https://github.com/mostafayasser)
- [Eslam Mamdouh \(https://github.com/EslamMamdoouh\)](https://github.com/EslamMamdoouh)

Problem Description and Background

Some people may not have the time to examine articles during their day because of time shortage.

Solution

The solution is to develop a learning agent to solve the problem. The agent will be a deep learning model that summarizes the output text using NLP techniques with Recurrent Neural Network to understand the context of the text and generate a summarization of the text.

Import required libraries

```
In [1]: import numpy as np
import pandas as pd
import re
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from nltk.corpus import stopwords
from tensorflow.keras.layers import Input, LSTM, Embedding, Dense, Concatenate
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.callbacks import EarlyStopping

import warnings
warnings.filterwarnings("ignore")
```

Using TensorFlow backend.

Load the dataset

```
In [2]: !wget --no-check-certificate https://storage.googleapis.com/islamohamed
```

```
data = pd.read_csv('data.csv', nrows=150000)
data.drop_duplicates(subset=['Text'], inplace=True)
data.dropna(axis=0, inplace=True)
```

```
--2020-05-13 08:17:47-- https://storage.googleapis.com/islamohamedd1.a
ppspot.com/amazon-fine-food-reviews/Reviews.csv (https://storage.google
apis.com/islamohamedd1.appspot.com/amazon-fine-food-reviews/Reviews.cs
v)
```

```
Resolving storage.googleapis.com (storage.googleapis.com)... 64.233.18
9.128, 2404:6800:4008:c07::80
```

```
Connecting to storage.googleapis.com (storage.googleapis.com)|64.233.18
9.128|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 300904694 (287M) [text/csv]
```

```
Saving to: 'data.csv'
```

```
data.csv          100%[=====>] 286.96M   195MB/s   in
1.5s
```

```
2020-05-13 08:17:49 (195 MB/s) - 'data.csv' saved [300904694/300904694]
```

Extract and format the training and testing data from the dataset

```
In [0]: # contraction_mapping source: https://gist.github.com/aravindpai/f21a286
contraction_mapping = {"ain't": "is not", "aren't": "are not", "can't": "
    "didn't": "did not", "doesn't": "does not", "
    "he'd": "he would", "he'll": "he will", "he's"
    "I'd": "I would", "I'd've": "I would have", "
    "i'd've": "i would have", "i'll": "i will",
    "it'd've": "it would have", "it'll": "it will
    "mayn't": "may not", "might've": "might have"
    "mustn't": "must not", "mustn't've": "must no
    "oughtn't": "ought not", "oughtn't've": "ough
    "she'd": "she would", "she'd've": "she would
    "should've": "should have", "shouldn't": "sho
    "this's": "this is", "that'd": "that would", "
    "there'd've": "there would have", "there's":
    "they'll": "they will", "they'll've": "they w
    "wasn't": "was not", "we'd": "we would", "we'
    "we've": "we have", "weren't": "were not", "w
    "what's": "what is", "what've": "what have",
    "where've": "where have", "who'll": "who will
    "why's": "why is", "why've": "why have", "wil
    "would've": "would have", "wouldn't": "would
    "y'all'd": "you all would", "y'all'd've": "you
    "you'd": "you would", "you'd've": "you would
    "you're": "you are", "you've": "you have"}

contraction_mapping['<br>'] = ""
```

```
In [4]: import nltk
from nltk.corpus import stopwords

nltk.download('stopwords')
nltk.download('punkt')

stop_words = set(stopwords.words('english'))

def clean_sentences(data):
    new_data = []
    for s in data:
        new_s = s.lower()
        new_s = ' '.join([w for w in s.split() if not w in stop_words])
        new_s = ' '.join([contraction_mapping[t] if t in contraction_mapping
                           else w for w, t in zip(s.split(), s.split())])
        new_data.append(new_s)
    return new_data

def clean_labels(data):
    new_data = []
    for s in data:
        new_s = s.lower()
        new_s = ' '.join([contraction_mapping[t] if t in contraction_mapping
                           else w for w, t in zip(s.split(), s.split())])
        new_data.append(new_s)
    return new_data

labels = clean_labels(data['Summary'])
sentences = clean_sentences(data['Text'])

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Unzipping tokenizers/punkt.zip.
```

```
In [0]: new_labels = []
for s in labels:
    new_s = s.split()
    new_s = ['_START_'] + new_s +['_END_']
    new_labels.append(' '.join(new_s))
labels = new_labels
```

```
In [6]: labels[:3]
```

```
Out[6]:['_START_ good quality dog food _END_',
'_START_ not as advertised _END_',
'_START_ "delight" says it all _END_']
```

```
In [0]: from sklearn.model_selection import train_test_split
training_sentences, testing_sentences, training_labels, testing_labels =
```

Tokenize the sentences

The data is converted to be numeric tokens so the model can work with it easily

```
In [0]: max_len = 80

x_tokenizer = Tokenizer()
x_tokenizer.fit_on_texts(training_sentences)

x_training = x_tokenizer.texts_to_sequences(training_sentences)
x_validation = x_tokenizer.texts_to_sequences(testing_sentences)

x_training = pad_sequences(x_training, maxlen=max_len, padding='post')
x_validation = pad_sequences(x_validation, maxlen=max_len, padding='post')

x_vocab_size = len(x_tokenizer.word_index) + 1
```

```
In [0]: y_tokenizer = Tokenizer()
y_tokenizer.fit_on_texts(training_labels)

y_training = y_tokenizer.texts_to_sequences(training_labels)
y_validation = y_tokenizer.texts_to_sequences(testing_labels)

y_training = pad_sequences(y_training, maxlen=max_len, padding='post')
y_validation = pad_sequences(y_validation, maxlen=max_len, padding='post')

y_vocab_size = len(y_tokenizer.word_index) + 1
```

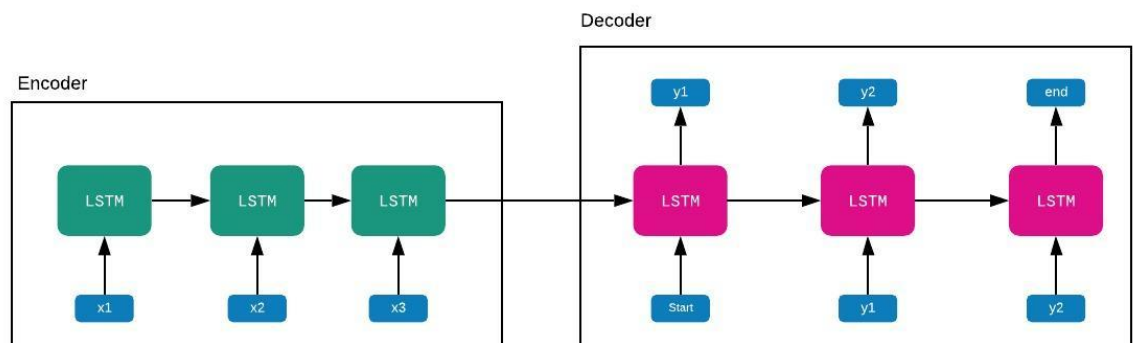
```
In [10]: y_tokenizer.word_counts['the']
```

```
Out[10]: 11208
```

Create the model archeticure

The model is using a LSTM (Long-Short-Term Memory) network for the encoder and another LSTM network for the decoder


The Encoder-Decoder archeticure will work as illustrated below



Encoder

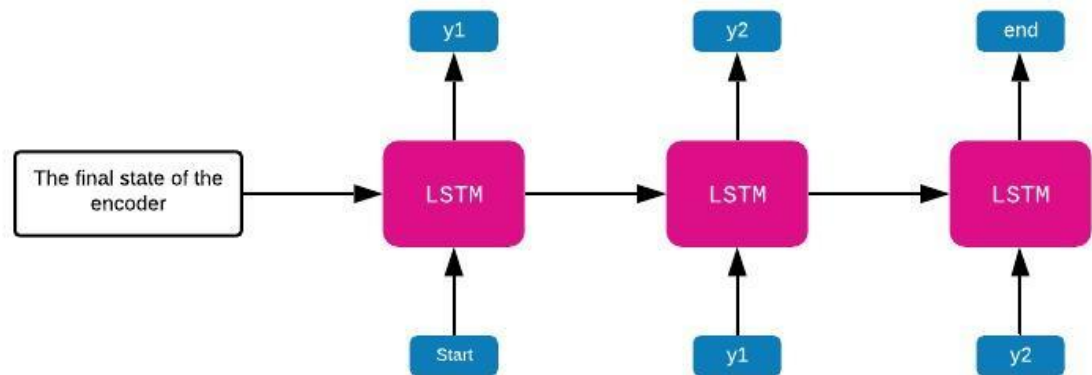
The encoder is composed of a LSTM network. The LSTM network starts with a 0 state as it's initial state. Then, the encoder network recieves a word every timestamp. the word is passed to the first LSTM layer and processed, the LSTM layer pass the state after processing the word to

the next LSTM layer. each LSTM layer receives the a word from the input sequence and the previous state. Finally, the encoder outputs the final state.



Decoder

The decoder LSTM network takes the final state of the encoder as it's initial state and the first word (start) as an input. Each LSTM layer in the decoder predicts the next word - y1, y2, etc - using the initial state of the encoder combined with the state of the previous LSTM state.



```

In [11]: from keras import backend as K
K.clear_session()
embedding_dim = 500

# Encoder
encoder_inputs = Input(shape=(max_len,))
encoder_embedding = Embedding(x_vocab_size, embedding_dim, trainable=True)

encoder_lstm1 = LSTM(embedding_dim, return_sequences=True, return_state=True)
encoder_output1, state_h1, state_c1 = encoder_lstm1(encoder_embedding)

encoder_lstm2 = LSTM(embedding_dim, return_sequences=True, return_state=True)
encoder_output2, state_h1, state_c1 = encoder_lstm2(encoder_output1)

encoder_lstm3 = LSTM(embedding_dim, return_state=True, return_sequences=True)
encoder_outputs, state_h, state_c = encoder_lstm3(encoder_output2)

# Decoder
decoder_inputs = Input(shape=(None,))
decoder_embedding_layer = Embedding(y_vocab_size, embedding_dim, trainable=True)
decoder_embedding = decoder_embedding_layer(decoder_inputs)

decoder_lstm = LSTM(embedding_dim, return_sequences=True, return_state=True)
decoder_outputs, decoder_fwd_state, decoder_back_state = decoder_lstm(decoder_embedding)

decoder_dense = TimeDistributed(Dense(y_vocab_size, activation='softmax'))
decoder_outputs = decoder_dense(decoder_outputs)

model = Model([encoder_inputs, decoder_inputs], decoder_outputs)
model.summary()

```

Model: "model"

Layer (type) connected to	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 80)]	0	

embedding (Embedding) t_1[0][0]	(None, 80, 500)	34941500	input_1

lstm (LSTM) dding[0][0]	[(None, 80, 500), (None, 200, 200)]	2002000	embedding

input_2 (InputLayer)	[(None, None)]	0	

lstm_1 (LSTM) [0][0]	[(None, 80, 500), (None, 200, 200)]	2002000	lstm

embedding_1 (Embedding) t_2[0][0]	(None, None, 500)	9138500	input
lstm_2 (LSTM) _1[0][0]	[(None, 80, 500), (N 2002000		lstm
lstm_3 (LSTM) dding_1[0][0]	[(None, None, 500), 2002000		embed
_2[0][1]			lstm
_2[0][2]			lstm
time_distributed (TimeDistribut _3[0][0]	(None, None, 18277)	9156777	lstm

=====

Total params: 61,244,777
Trainable params: 61,244,777
Non-trainable params: 0

=====

```
In [0]: model.compile(optimizer='rmsprop', loss='sparse_categorical_crossentropy
```

Create an early stopping callback function

A callback function is called after each training epoch and stops the training of the validation loss started to increase

```
In [0]: early_stopping_callback = EarlyStopping(monitor='val_loss', mode='min',
```

Mount google drive to the file system (Only in google colab)


```
In [14]: from google.colab import drive
drive.mount('/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly (https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly)

Enter your authorization code:

.....

Mounted at /drive

Create a callback function to save the weights of the model after each epoch

```
In [0]: import os
import tensorflow as tf

checkpoint_path = "/drive/My Drive/ai_project/reviews/text_sum_model_weights"
model_path = "/drive/My Drive/ai_project/reviews/text_sum_model"
checkpoint_dir = os.path.dirname(checkpoint_path)
```

```
In [0]: # Create a callback that saves the model's weights
checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path,
                                                         save_weights_only=True,
                                                         verbose=1)
```

Train the model

```
In [17]: history=model.fit(
        [x_training,y_training[:, :-1]],
        y_training.reshape(y_training.shape[0],y_training.shape[1], 1)[:,1:],
        epochs=50,
        callbacks=[early_stopping_callback, checkpoint_callback],
        batch_size=512,
        verbose=1,
        validation_data=(x_validation,y_validation[:, :-1]), y_validation.re
    )
model.save(model_path)
```

Epoch 1/50

225/225 [=====] - ETA: 0s - loss: 0.5200 - acc: 0.9419

Epoch 00001: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt

225/225 [=====] - 262s 1s/step - loss: 0.5200 - acc: 0.9419 - val_loss: 0.3448 - val_acc: 0.9501

Epoch 2/50

225/225 [=====] - ETA: 0s - loss: 0.3375 - acc: 0.9510

Epoch 00002: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt

225/225 [=====] - 260s 1s/step - loss: 0.3375 - acc: 0.9510 - val_loss: 0.3152 - val_acc: 0.9527

Epoch 3/50

225/225 [=====] - ETA: 0s - loss: 0.3151 - acc: 0.9529

Epoch 00003: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt

225/225 [=====] - 260s 1s/step - loss: 0.3151 - acc: 0.9529 - val_loss: 0.3015 - val_acc: 0.9540

Epoch 4/50

225/225 [=====] - ETA: 0s - loss: 0.3002 - acc: 0.9540

Epoch 00004: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt

225/225 [=====] - 261s 1s/step - loss: 0.3002 - acc: 0.9540 - val_loss: 0.2896 - val_acc: 0.9548

Epoch 5/50

225/225 [=====] - ETA: 0s - loss: 0.2877 - acc: 0.9548

Epoch 00005: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt

225/225 [=====] - 261s 1s/step - loss: 0.2877 - acc: 0.9548 - val_loss: 0.2831 - val_acc: 0.9553

Epoch 6/50

225/225 [=====] - ETA: 0s - loss: 0.2774 - acc: 0.9555

Epoch 00006: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt

225/225 [=====] - 260s 1s/step - loss: 0.2774 - acc: 0.9555 - val_loss: 0.2776 - val_acc: 0.9556

Epoch 7/50

225/225 [=====] - ETA: 0s - loss: 0.2680 - acc: 0.9563



Epoch 00007: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 261s 1s/step - loss: 0.2680
- acc: 0.9563 - val_loss: 0.2725 - val_acc: 0.9563
Epoch 8/50
225/225 [=====] - ETA: 0s - loss: 0.2595 - acc: 0.9570
Epoch 00008: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 260s 1s/step - loss: 0.2595
- acc: 0.9570 - val_loss: 0.2681 - val_acc: 0.9566
Epoch 9/50
225/225 [=====] - ETA: 0s - loss: 0.2518 - acc: 0.9576
Epoch 00009: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 260s 1s/step - loss: 0.2518
- acc: 0.9576 - val_loss: 0.2652 - val_acc: 0.9569
Epoch 10/50
225/225 [=====] - ETA: 0s - loss: 0.2446 - acc: 0.9582
Epoch 00010: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 261s 1s/step - loss: 0.2446
- acc: 0.9582 - val_loss: 0.2630 - val_acc: 0.9571
Epoch 11/50
225/225 [=====] - ETA: 0s - loss: 0.2375 - acc: 0.9588
Epoch 00011: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 259s 1s/step - loss: 0.2375
- acc: 0.9588 - val_loss: 0.2624 - val_acc: 0.9571
Epoch 12/50
225/225 [=====] - ETA: 0s - loss: 0.2307 - acc: 0.9594
Epoch 00012: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 261s 1s/step - loss: 0.2307
- acc: 0.9594 - val_loss: 0.2610 - val_acc: 0.9573
Epoch 13/50
225/225 [=====] - ETA: 0s - loss: 0.2241 - acc: 0.9600
Epoch 00013: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 260s 1s/step - loss: 0.2241
- acc: 0.9600 - val_loss: 0.2609 - val_acc: 0.9573
Epoch 14/50
225/225 [=====] - ETA: 0s - loss: 0.2174 - acc: 0.9607
Epoch 00014: saving model to /drive/My Drive/ai_project/reviews/text_sum_model_weights/cp.ckpt
225/225 [=====] - 260s 1s/step - loss: 0.2174
- acc: 0.9607 - val_loss: 0.2614 - val_acc: 0.9572
Epoch 00014: early stopping
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/resource_variable_ops.py:1817: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with con

straint is deprecated and will be removed in a future version.
Instructions for updating:
If using Keras pass *_constraint arguments to layers.
INFO:tensorflow:Assets written to: /drive/My Drive/ai_project/reviews/t
ext_sum_model/assets

Plot the progress of the loss and the accuracy during the training

```
In [18]: from matplotlib import pyplot
pyplot.plot(history.history['loss'], label='train loss')
pyplot.plot(history.history['val_loss'], label='test loss')
pyplot.legend()
pyplot.show()

pyplot.plot(history.history['acc'], label='train accuracy')
pyplot.plot(history.history['val_acc'], label='test accuracy')
pyplot.legend()
pyplot.show()
```



```
In [0]: # model.load_weights(checkpoint_path)
```

Create a reversed word_index to be able to get the text of each word

using it's token

```
In [0]: reverse_target_word_index=y_tokenizer.index_word
reverse_source_word_index=x_tokenizer.index_word
target_word_index=y_tokenizer.word_index
```

```
In [0]: # encoder inference
encoder_model = Model(inputs=encoder_inputs,outputs=[encoder_outputs, st

#decoder inference
# Below tensors will hold the states of the previous time step
decoder_state_input_h = Input(shape=(embedding_dim,))
decoder_state_input_c = Input(shape=(embedding_dim,))
decoder_hidden_state_input = Input(shape=(max_len, embedding_dim))

# Get the embeddings of the decoder sequence
decoder_embedding2 = decoder_embedding_layer(decoder_inputs)

# To predict the next word in the sequence, set the initial states to th
decoder_outputs2, state_h2, state_c2 = decoder_lstm(decoder_embedding2,

decoder_outputs2 = decoder_dense(decoder_outputs2)

decoder_model = Model(
    [decoder_inputs] + [decoder_hidden_state_input,decoder_state_input_h
    [decoder_outputs2] + [state_h2, state_c2]
)
```

```

In [0]: def decode_sequence(input_sequence):
        encoder_out, encoder_h, encoder_c = encoder_model.predict(input_sequ

        output_sequence = np.zeros((1,1))

        output_sequence[0, 0] = target_word_index['start']

        stop = False
        decoded_sentence = ''
        while not stop:
            output_tokens, h, c = decoder_model.predict([output_sequence] +

            heighest_probability_token = np.argmax(output_tokens[0, -1, :])
            if (heighest_probability_token == 0):
                break
            hieighest_probability_word = reverse_target_word_index[heighest_p

            if (hieighest_probability_word != 'end'):
                decoded_sentence += ' ' + hieighest_probability_word

            if (hieighest_probability_word == 'end' or len(decoded_sentence.s
                stop = True

            # Update the target sequence (of length 1).
            output_sequence = np.zeros((1,1))
            output_sequence[0, 0] = heighest_probability_token

            # Update internal states
            encoder_h, encoder_c = h, c

        return decoded_sentence

```

Test the model

```

In [0]: def get_summary(input_sequence):
        newString=''
        for i in input_sequence:
            if((i!=0 and i!=target_word_index['start']) and i!=target_word_ind
                newString=newString+reverse_target_word_index[i]+' '
        return newString

        def get_text(input_sequence):
            newString = ''
            for i in input_sequence:
                if(i != 0):
                    newString = newString + reverse_source_word_index[i] + ' '
            return newString

```

```
In [24]: for i in range(200):  
         original_summary = get_summary(y_validation[i])  
         predicted_summary = decode_sequence(x_validation[i].reshape(1,max_len))  
         print("Review:", get_text(x_validation[i]))  
         print("Original summary:", original_summary)  
         print("Predicted summary:", predicted_summary)  
  
         print("\n")
```

y science diet s fish flavored food give try

Original summary: not tasty

Predicted summary: my cats love it

Review: my cats agree really love dry food i surprised even seemed like better epigen 90 i normally feed them no barf potty issues either eating food week whew lets hope still case i try non hairball formula

Original summary: chicken is paw good

Predicted summary: my cats love it

Review: works well power breeds heavy chewers if cheaper id buy dozen we pitbulls rottweilers far best choice keeping busy

Original summary: great for larger dogs

Predicted summary: great product

Review: this negative 3 star review the noodles squid ish taste texture i nearly gagged trying are bad or no calorie mush plant solubles r

In [0]: